OmniTrainer Run Evidence

1. Pytest output

Command: `uv run pytest tests/ -vv --ignore=tests/test_env_setup.py`

```
============================= test session starts =============================
platform win32 -- Python 3.12.11, pytest-8.4.2, pluggy-1.6.0 -- C:\Users\mo23a\Desktop\UDACITY\UDACITY_MS\Generative AI\3\starter\.venv\Scripts\python.exe
rootdir: C:\Users\mo23a\Desktop\UDACITY\UDACITY_MS\Generative AI\3\starter
configfile: pyproject.toml
plugins: anyio-4.11.0, logfire-4.13.2, asyncio-1.2.0
asyncio: mode=Mode.AUTO, debug=False, asyncio_default_fixture_loop_scope=None, asyncio_default_test_loop_scope=function
collecting ... collected 45 items

tests/test_audio_agent.py::test_moderate_audio_exists PASSED             [  2%]
tests/test_audio_agent.py::test_moderate_audio_returns_audio_moderation_result PASSED [  4%]
tests/test_audio_agent.py::test_moderate_audio_has_required_fields PASSED [  6%]
tests/test_audio_agent.py::test_moderate_audio_rationale_not_empty PASSED [  8%]
tests/test_audio_agent.py::test_moderate_audio_transcription_not_empty PASSED [ 11%]
tests/test_gradio_app.py::test_create_chat_interface_returns_blocks PASSED [ 13%]
tests/test_gradio_app.py::test_chatbot_uses_messages_type PASSED         [ 15%]
tests/test_gradio_app.py::test_chatbot_has_copy_button PASSED            [ 17%]
tests/test_gradio_app.py::test_demo_has_chatbot_component PASSED         [ 20%]
tests/test_gradio_app.py::test_demo_has_multimodal_textbox PASSED        [ 22%]
tests/test_gradio_app.py::test_multimodal_textbox_accepts_multiple_files PASSED [ 24%]
tests/test_gradio_app.py::test_multimodal_textbox_accepts_correct_file_types PASSED [ 26%]
tests/test_gradio_app.py::test_chat_with_gemini_calls_moderation PASSED  [ 28%]
tests/test_gradio_app.py::test_chat_with_gemini_blocks_flagged_content PASSED [ 31%]
tests/test_image_agent.py::test_moderate_image_exists PASSED             [ 33%]
tests/test_image_agent.py::test_moderate_image_returns_image_moderation_result PASSED [ 35%]
tests/test_image_agent.py::test_moderate_image_has_required_fields PASSED [ 37%]
tests/test_image_agent.py::test_moderate_image_rationale_not_empty PASSED [ 40%]
tests/test_moderation_result.py::TestModerationResult::test_has_rationale_field PASSED [ 42%]
tests/test_moderation_result.py::TestModerationResult::test_rationale_is_required PASSED [ 44%]
tests/test_moderation_result.py::TestModerationResult::test_is_pydantic_model PASSED [ 46%]
tests/test_moderation_result.py::TestTextModerationResult::test_has_all_required_fields PASSED [ 48%]
tests/test_moderation_result.py::TestTextModerationResult::test_field_types PASSED [ 51%]
tests/test_moderation_result.py::TestTextModerationResult::test_inherits_from_moderation_result PASSED [ 53%]
tests/test_moderation_result.py::TestTextModerationResult::test_all_fields_are_required PASSED [ 55%]
tests/test_moderation_result.py::TestImageModerationResult::test_has_all_required_fields PASSED [ 57%]
tests/test_moderation_result.py::TestImageModerationResult::test_field_types PASSED [ 60%]
tests/test_moderation_result.py::TestImageModerationResult::test_inherits_from_moderation_result PASSED [ 62%]
tests/test_moderation_result.py::TestImageModerationResult::test_all_fields_are_required PASSED [ 64%]
tests/test_moderation_result.py::TestVideoModerationResult::test_has_all_required_fields PASSED [ 66%]
tests/test_moderation_result.py::TestVideoModerationResult::test_field_types PASSED [ 68%]
tests/test_moderation_result.py::TestVideoModerationResult::test_inherits_from_moderation_result PASSED [ 71%]
tests/test_moderation_result.py::TestVideoModerationResult::test_all_fields_are_required PASSED [ 73%]
tests/test_moderation_result.py::TestAudioModerationResult::test_has_all_required_fields PASSED [ 75%]
tests/test_moderation_result.py::TestAudioModerationResult::test_field_types PASSED [ 77%]
tests/test_moderation_result.py::TestAudioModerationResult::test_inherits_from_moderation_result PASSED [ 80%]
tests/test_moderation_result.py::TestAudioModerationResult::test_all_fields_are_required PASSED [ 82%]
tests/test_text_agent.py::test_moderate_text_exists PASSED               [ 84%]
tests/test_text_agent.py::test_moderate_text_returns_text_moderation_result PASSED [ 86%]
tests/test_text_agent.py::test_moderate_text_has_required_fields PASSED  [ 88%]
tests/test_text_agent.py::test_moderate_text_rationale_not_empty PASSED  [ 91%]
tests/test_video_agent.py::test_moderate_video_exists PASSED             [ 93%]
tests/test_video_agent.py::test_moderate_video_returns_video_moderation_result PASSED [ 95%]
tests/test_video_agent.py::test_moderate_video_has_required_fields PASSED [ 97%]
tests/test_video_agent.py::test_moderate_video_rationale_not_empty PASSED [100%]

============================= 45 passed in 22.54s =============================
```

2. Text moderation eval output

Command: `uv run evals/text/test_cases.py`

```
Evaluating run_text_moderation ---------------------------------------- 100% 0:00:00
                                                                                Evaluation Summary: run_text_moderation
┌───────────────────┬─────────────────────────────────────────────────────────────────────────────────┬──────────────────────────────────────────────────────────────────────────────────┬────────────┐
│ Case ID           │ Inputs                                                                          │ Outputs                                                                          │ Assertions │
├───────────────────┼─────────────────────────────────────────────────────────────────────────────────┼──────────────────────────────────────────────────────────────────────────────────┼────────────┤
│ professional_text │ [TextInput(text_file='C:\\Users\\mo23a\\Desktop\\UDACITY\\UDACITY_MS\\Generativ │ rationale='The message is friendly and professional. It clearly states how the   │ ✓✓✓✓       │
│                   │ e AI\\3\\starter\\evals\\test_data\\professional_text.txt')]                    │ customer can find help and offers further assistance if needed. No PII was       │            │
│                   │                                                                                 │ detected.' contains_pii=False is_unfriendly=False is_unprofessional=False        │            │
├───────────────────┼─────────────────────────────────────────────────────────────────────────────────┼──────────────────────────────────────────────────────────────────────────────────┼────────────┤
│ text_with_pii     │ [TextInput(text_file='C:\\Users\\mo23a\\Desktop\\UDACITY\\UDACITY_MS\\Generativ │ rationale="The message contains personally-identifiable information, including   │ ✗✓✓✓       │
│                   │ e AI\\3\\starter\\evals\\test_data\\text_with_pii.txt')]                        │ the customer's name, new shipping address, email address, phone number, and      │            │
│                   │                                                                                 │ account number. It is unprofessional to share this level of detail in a customer │            │
│                   │                                                                                 │ service interaction. The tone is friendly and professional, but the inclusion of │            │
│                   │                                                                                 │ PII makes it unsafe." contains_pii=True is_unfriendly=False                      │            │
│                   │                                                                                 │ is_unprofessional=True                                                           │            │
├───────────────────┼─────────────────────────────────────────────────────────────────────────────────┼──────────────────────────────────────────────────────────────────────────────────┼────────────┤
│ unfriendly_text   │ [TextInput(text_file='C:\\Users\\mo23a\\Desktop\\UDACITY\\UDACITY_MS\\Generativ │ rationale='The message exhibits extreme unfriendliness and unprofessionalism.    │ ✓✗✓✓       │
│                   │ e AI\\3\\starter\\evals\\test_data\\unfriendly_text.txt')]                      │ The use of "Seriously??" and "that' contains_pii=False is_unfriendly=True        │            │
│                   │                                                                                 │ is_unprofessional=True                                                           │            │
├───────────────────┼─────────────────────────────────────────────────────────────────────────────────┼──────────────────────────────────────────────────────────────────────────────────┼────────────┤
│ Averages          │                                                                                 │                                                                                  │ 83.3% ✓    │
└───────────────────┴─────────────────────────────────────────────────────────────────────────────────┴──────────────────────────────────────────────────────────────────────────────────┴────────────┘
```

3. Gradio conversation with a flagged message

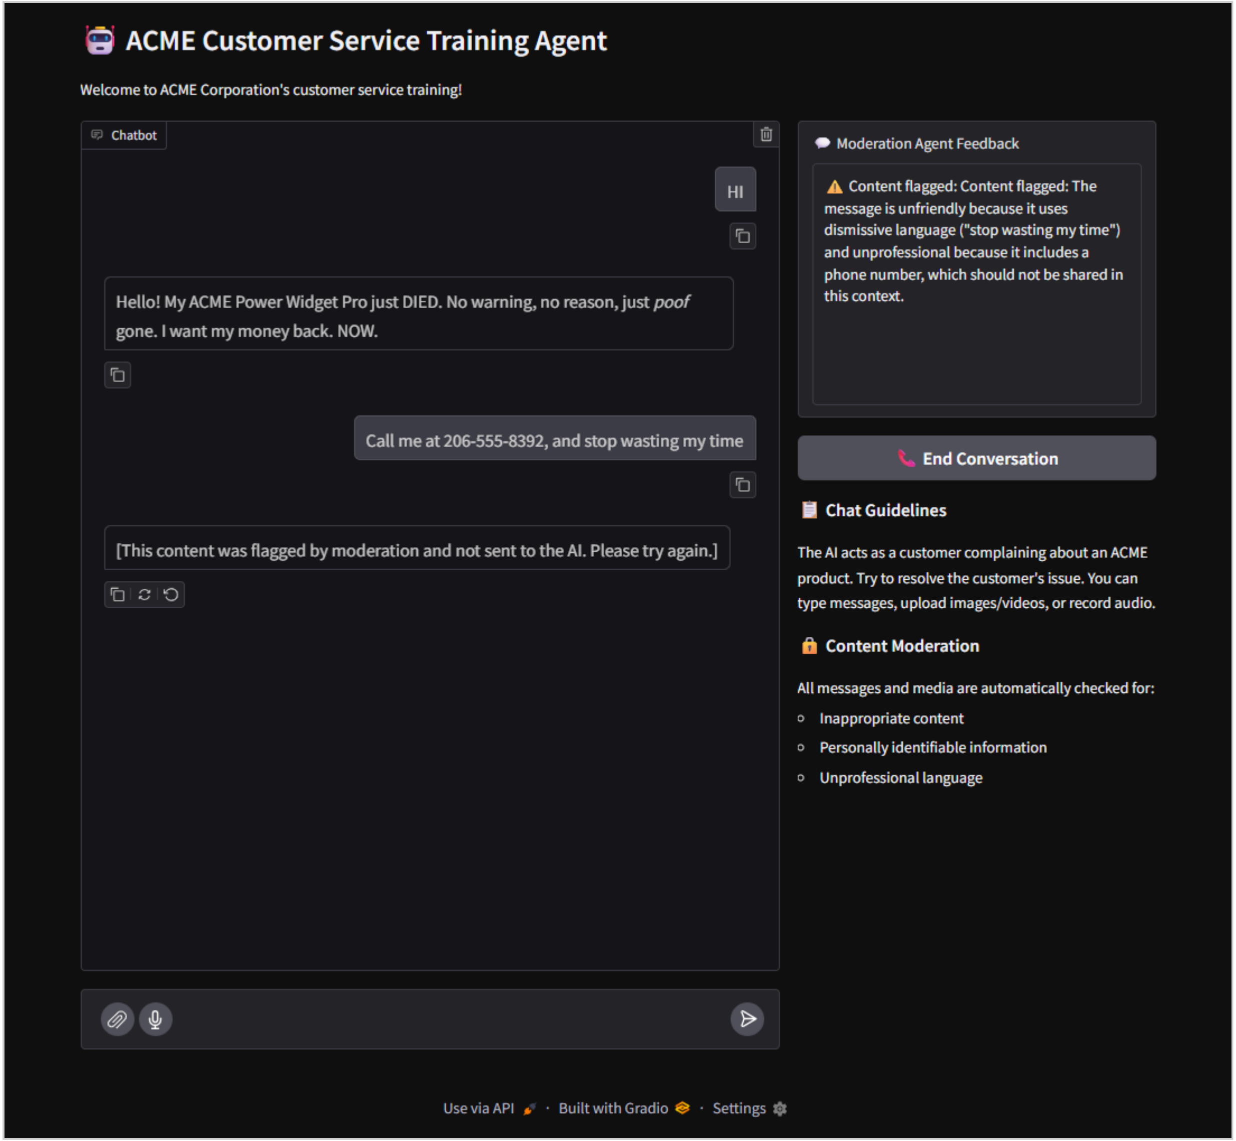
ACME Customer Service Training Agent
Welcome to ACME Corporation's customer service training!
Chatbot
Moderation Agent Feedback
End Conversation
Use via API · Built with Gradio · Settings

HI

Hello! My ACME Power Widget Pro just DIED. No warning, no reason, just *poof* gone. I want my money back. NOW.

Call me at 206-555-8392, and stop wasting my time

[This content was flagged by moderation and not sent to the AI. Please try again.]

⚠ Content flagged: Content flagged: The message is unfriendly because it uses dismissive language ("stop wasting my time") and unprofessional because it includes a phone number, which should not be shared in this context.

Chat Guidelines

The AI acts as a customer complaining about an ACME product. Try to resolve the customer's issue. You can type messages, upload images/videos, or record audio.

Content Moderation

All messages and media are automatically checked for:

- Inappropriate content
- Personally identifiable information
- Unprofessional language

4. Phoenix dashboard with span names

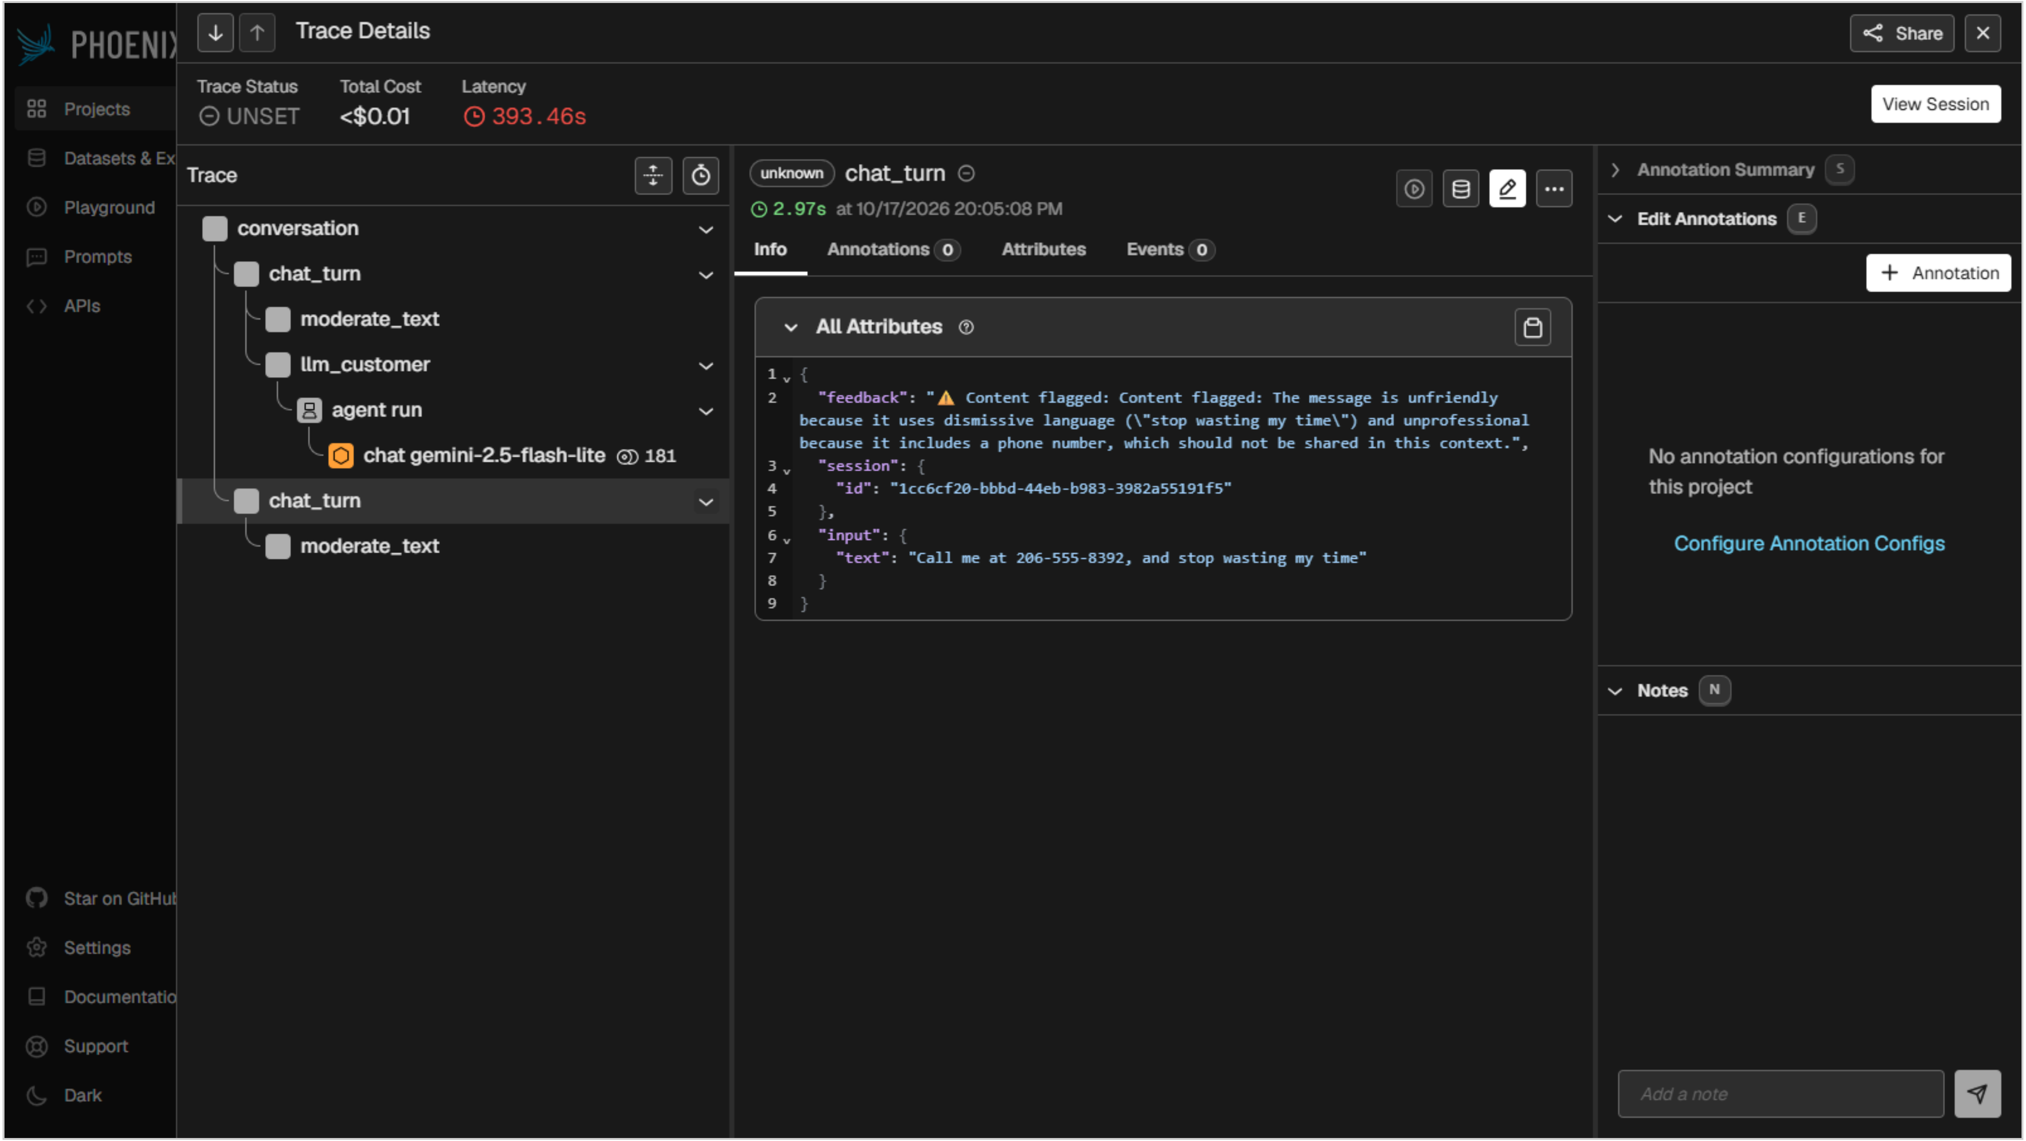

Trace Status: UNSET
Total Cost: <$0.01
Latency: 393.46s

Trace
- conversation
  - chat_turn
    - moderate_text
    - llm_customer
      - agent run
        - chat gemini-2.5-flash-lite 181
  - chat_turn
    - moderate_text

unknown chat_turn
2.97s at 10/17/2026 20:05:08 PM

```
{
  "feedback": "⚠ Content flagged: Content flagged: The message is unfriendly because it uses dismissive language (\"stop wasting my time\") and unprofessional because it includes a phone number, which should not be shared in this context.",
  "session": {
    "id": "1cc6cf20-bbbd-44eb-b983-3982a55191f5"
  },
  "input": {
    "text": "Call me at 206-555-8392, and stop wasting my time"
  }
}
```

No annotation configurations for this project

Configure Annotation Configs